

Enterprise Agentic Operations Platform

A secure, evaluable multi-agent system for enterprise knowledge, investigation, and controlled action.

Hybrid RAG

MCP tools

Human approval

Security gates



THE PROBLEM

Enterprise AI needs more than a chat interface

Useful automation must be grounded in authorized data, safe around sensitive information, observable in production, and accountable when it takes action.

01

Grounded knowledge

Hybrid retrieval, reranking, citations, and abstention when evidence is insufficient.

02

Controlled actions

MCP tools with typed contracts, least privilege, retries, and human approval for writes.


03

Operational trust

PII masking, tenant isolation, audit trails, evaluations, cost budgets, and tracing.

"An incident-response workflow that can investigate, explain, and act—without silently crossing a security boundary."

One platform, replaceable capabilities



Incident response from question to approved action

1 · Secure intake

Authenticate user, resolve tenant and roles, scan input for injection and PII.

2 · Investigate

Retrieve authorized evidence, rerank sources, and call read-only status tools.

3 · Explain

Return a concise diagnosis with citations, confidence, provenance, and next steps.

4 · Act safely

Request approval, validate scope, create a ticket, and record the complete audit event.



0

unauthorized sources in context

100%

write actions require approval

1

auditable workflow trace

Security is a system boundary, not a prompt

Before the model

- Authentication and tenant isolation
- Role-aware retrieval filters
- PII detection and masking
- Prompt-injection scanning

Before any tool action

- Strict typed schemas
- Allowlisted MCP tools
- Read versus write permissions
- Human approval for side effects

After model output

- Citation and claim validation
- Secret and PII scanning
- Safe refusal and abstention
- Response schema enforcement

Always observable

- Prompt and retrieval versions
- Tool arguments and results
- Latency, tokens, and cost
- Immutable audit events

Least privilege applies to data, tools, model context, and the final action.

Every improvement is measurable

R

Retrieval quality

Recall@K, MRR, NDCG, ACL correctness, citation completeness, and latency.

A

Agent quality

Groundedness, answer correctness, tool selection, recovery, and appropriate abstention.

S

Security quality

Injection resistance, PII leakage, cross-tenant access, and unauthorized-action tests.

Scale strategy

Async ingestion and workflows, Redis queues, caching, model routing, rate limits, and per-workflow budgets.

Provider flexibility

Model, embedding, retrieval, MCP, and observability providers sit behind typed interfaces and can be benchmarked independently.

The platform earns trust through evidence, constrained actions, and regression-tested behavior.